

Amazon EC2 On-Demand Pricing Notes (us-east-1, sample reference data for this exercise)

Instance Type | vCPU | Memory (GiB) | On-Demand Price (USD/hour)

t3.medium | 2 | 4 | 0.0416

t3.large | 2 | 8 | 0.0832

m5.large | 2 | 8 | 0.096

m5.xlarge | 4 | 16 | 0.192

r5.large | 2 | 16 | 0.126

c5.large | 2 | 4 | 0.085

General guidance:

- On-Demand pricing is billed per second, with a one-minute minimum, and no long-term commitment.
- Typical well-optimized EC2 workloads for small-to-medium services fall in the range of \$0.04 to \$0.20 per hour, depending on instance family and size.
- An hourly rate significantly above \$1.00/hour for a single instance usually indicates either a much larger instance size (many vCPUs / high memory), a GPU-accelerated instance type, or a misconfigured / oversized resource that should be reviewed.
- Reserved Instances and Savings Plans can reduce On-Demand rates by 30-72% in exchange for a 1 or 3 year commitment.
- Spot Instances can reduce cost by up to 90% versus On-Demand, but can be interrupted with short notice and are not suitable for workloads that cannot tolerate interruption.
- Data transfer OUT to the internet is billed separately from compute hours, and is not included in the hourly instance rate.